

# Sagar Malhotra

Machine Learning Research Unit,  
TU Wien, Austria

Email: [sagar.malhotra@tuwien.ac.at](mailto:sagar.malhotra@tuwien.ac.at)  
Personal website: [countinglogic.github.io](https://countinglogic.github.io)

## Academic Employment

2023-now    **Postdoctoral Researcher**  
Host: Prof. Thomas Gärtner  
Machine Learning Research Unit,  
TU Wien (Technical University of Vienna), Austria

## Education

2019-2023    **PhD in Computer Science**  
Thesis: On Tractability and Consistency of Probabilistic Inference in Relational Domains  
Advisor: Prof. Luciano Serafini  
University of Trento, Italy  
Fondazione Bruno Kessler, Italy (dual affiliation)

2016-2018    **MSc in Physics**  
University of Trento, Italy

2012-2015    **BSc in Physics (Honors)**  
University of Delhi, India

## Research Interests

My research focuses on algorithmic foundations of sound, efficient, and safe machine learning. In other words, my goal is to develop practical machine learning algorithms with formal guarantees on the aforementioned properties. I am particularly interested in problems at the intersection of logic, probability and machine learning. My interests and expertise have allowed me to work on: tractability of learning and reasoning over relational data; logical expressivity based explainability methods; and quantitative verification of machine learning algorithms.

## Conference Publications<sup>\*</sup>

2024    Alexander Pluska, Pascal Welke, Thomas Gärtner and **Sagar Malhotra**.  
Logical Distillation of Graph Neural Networks  
*International Conference on Principles of Knowledge Representation and Reasoning 2024*  
[KR 2024](#) (CORE Rank A\*, 17% acceptance rate in the special track)

---

<sup>\*</sup>Supervised student coauthors are underlined.

- 2024 Florian Chen, Felix Weitkämper, and **Sagar Malhotra**.  
Understanding Domain-Size Generalization in Markov Logic Networks  
*Machine Learning and Knowledge Discovery in Databases. Research Track - European Conference, ECML PKDD 2024*  
[ECML PKDD 2024](#) (CORE Rank A, 24% acceptance rate)
- 2024 Alessandro Daniele, Tommaso Campari, **Sagar Malhotra** and Luciano Serafini  
Simple and Effective Transfer Learning for Neuro-Symbolic Integration  
International Conference on Neural-Symbolic Learning and Reasoning, NeSy 2024  
[NeSy 2024](#) (Best Paper Award)
- 2023 Alessandro Daniele, Tommaso Campari, **Sagar Malhotra** and Luciano Serafini.  
Deep Symbolic Learning: Discovering Symbols and Rules from Perception  
*International Joint Conference on Artificial Intelligence 2023*  
[IJCAI 2023](#) (CORE Rank A\*, 15% acceptance rate)
- 2022 **Sagar Malhotra** and Luciano Serafini  
On Projectivity in Markov Logic Networks  
*Machine Learning and Knowledge Discovery in Databases. Research Track - European Conference, ECML PKDD 2022*  
[ECML PKDD 2022](#)(CORE Rank A, 26% acceptance rate).
- 2022 **Sagar Malhotra** and Luciano Serafini  
Weighted Model Counting in  $FO^2$  with Cardinality Constraints and Counting Quantifiers:  
A Closed Form Formula  
*AAAI Conference on Artificial Intelligence 2022*  
[AAAI 2022](#) (CORE Rank A\*, 15% acceptance rate, accepted as oral presentation)
- 2021 **Sagar Malhotra** and Luciano Serafini  
A Combinatorial Approach to Weighted Model Counting in the Two Variable Fragment  
with Cardinality Constraints  
*International Conference of the Italian Association for Artificial Intelligence 2019*  
[AIxIA 2021](#)

## Workshop Publications<sup>\*</sup>

- 2024 Patrick Indri, Peter Blohm, Anagha Athavale, Ezio Bartocci, Georg Weissenbacher, Matteo Maffei, Dejan Nickovic, Thomas Gärtner, **Sagar Malhotra**  
Distillation based Robustness Verification with PAC Guarantees  
*Next Generation of AI Safety Workshop, ICML 2024*  
[NextGenAISafety, ICML 2024](#)
- 2024 Alexander Pluska, Pascal Welke, Thomas Gärtner and **Sagar Malhotra**.  
Logical Distillation of Graph Neural Networks  
*Workshop on Mechanistic Interpretability, ICML 2024*  
[MI Workshop, ICML 2024](#)

---

<sup>\*</sup>Supervised student coauthors are underlined

- 2023 Alessandro Daniele, Tommaso Campari, **Sagar Malhotra** and Luciano Serafini.  
Deep Symbolic Learning: Discovering Symbols and Rules from Perception  
*International Workshop on Neural-Symbolic Learning and Reasoning 2023*  
[NeSy 2023](#) (Accepted for spotlight presentation)
- 2022 **Sagar Malhotra** and Luciano Serafini  
On Projectivity in Markov Logic Networks  
*International Workshop on Probabilistic Logic Programming 2022*  
[PLP 2022](#)
- 2021 **Sagar Malhotra** and Luciano Serafini. Weighted Model Counting in  $FO^2$  with Cardinality Constraints and Counting Quantifiers: A Closed Form Formula  
*International Workshop on Statistical Relational AI, IJCLR 2021*.  
[StarAI, IJCLR 2021](#)
- 2020 **Sagar Malhotra** and Luciano Serafini. Weighted Model Counting in  $C^2$  (Abstract)  
*Workshop on Machine Learning and Data Mining, AIXIA 2020*  
[MLDM 2020](#)

## Preprints<sup>\*</sup>

- 2024 Davide Bizzaro, Luciano Serafini and **Sagar Malhotra**  
Towards Counting Markov Equivalence Classes with Logical Constraints  
[Arxiv](#)
- 2023 **Sagar Malhotra**, Davide Bizzaro and Luciano Serafini  
Lifted Inference beyond First Order Logic  
*Under Review at Artificial Intelligence Journal (Major Revision)*  
[Arxiv](#)

## Talks and Tutorials

- 2024 Fundamental Problems in Statistical Relational AI  
Tutorial proposal accepted for KR 2024
- 2022 On Consistency of Learning and Inference in Statistical Relational Learning  
*Invited Talk at MLDM Workshop at the AIXIA Conference 2022, Udine, Italy* ([Abstract](#))
- 2022 On Probabilistic Inference in Logical Domains  
*Invited talk at the Institute of Informatics, Ludwig Maximilian University of Munich, Germany*
- 2022 A Tutorial on Probabilistic Inference in Logical Domains  
*Guest Lecture at the Knowledge representation and Learning course, University of Padova, Italy*
- 2022 Weighted First-Order Model Counting  
*DocInProgress Colloquium, Department of Mathematics, University of Trento, Italy*

---

<sup>\*</sup>Supervised students coauthors are underlined

2022      Weighted First-Order Model Counting  
AAAI 2022@FBK Workshop, Trento, Italy ([Video](#))

## Reviewing and PC Experience

Reviewer ICLR 2024, PC member AAAI 2024, Reviewer NeurIPS 2024, Reviewer KR 2024, PC Member IJCAI 2024, PC Member PLP workshop 2024, Reviewer ICLR 2024, Reviewer AISTATS 2024, PC Member AAAI 2024, PC Member SAC 2024, Reviewer NeurIPS 2023, PC Member MLG Workshop-ECML 2023, PC Member PLP workshop 2023, PC Member KR 2023, PC Member AAAI 2023, Reviewer AISTATS 2023, Sub-Reviewer KR 2021, Reviewer for Data Mining and Knowledge Discovery (Q1 Journal)

## Master's Student Supervision (See Teaching Statement )

2024      Peter Blohm, TU Wien, Austria  
Thesis title (tentative): Practical PAC-Verification with Signal Temporal Logic

2023      Davide Bizzaro, University of Padova, Italy  
Thesis: Lifted Inference Beyond First Order Logic

## Other important Student Supervision Roles

- Florian Chen, TU Wien, Austria  
(Bachelor's Intern at the Research Unit, leading to a conference publication)
- Alexander Pluska, TU Wien, Austria  
(Supervised in a graduate course, leading to a conference publication)

## Teaching Experience (See Teaching Statement)

2024  
(Winter Semester)      **Modern Applications of Logic in Machine Learning**  
(Elective graduate course in Computer Science, dean's approval awaited)

2023 - Now      **Introduction to Machine Learning** (Bachelor's in Computer Science)

2023 - Now      **Machine Learning Algorithms and Applications** (Master's in Computer Science)

2023 - Now      **Theoretical Foundations and Research Topics in Machine Learning**  
(Master's in Computer Science)